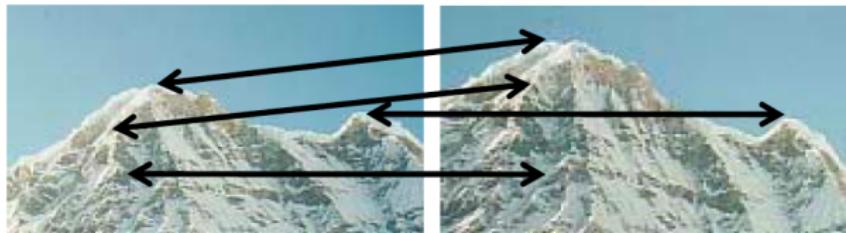
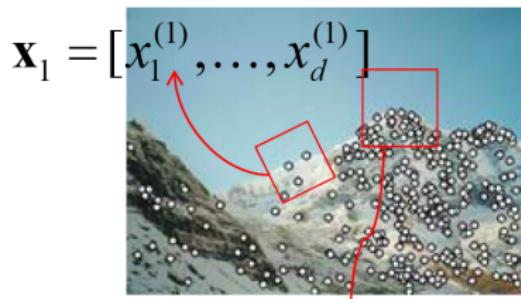


Image Features: Local Descriptors

Local Features

- **Detection:** Identify the interest points.
- **Description:** Extract a feature descriptor around each interest point.
- **Matching:** Determine correspondence between descriptors in two views.



[Source: K. Grauman]

The Ideal Feature Descriptor

- **Repeatable:** Invariant to rotation, scale, photometric variations
- **Distinctive:** We will need to match it to lots of images/objects!
- **Compact:** Should capture rich information yet not be too high-dimensional (otherwise matching will be slow)
- **Efficient:** We would like to compute it (close-to) real-time

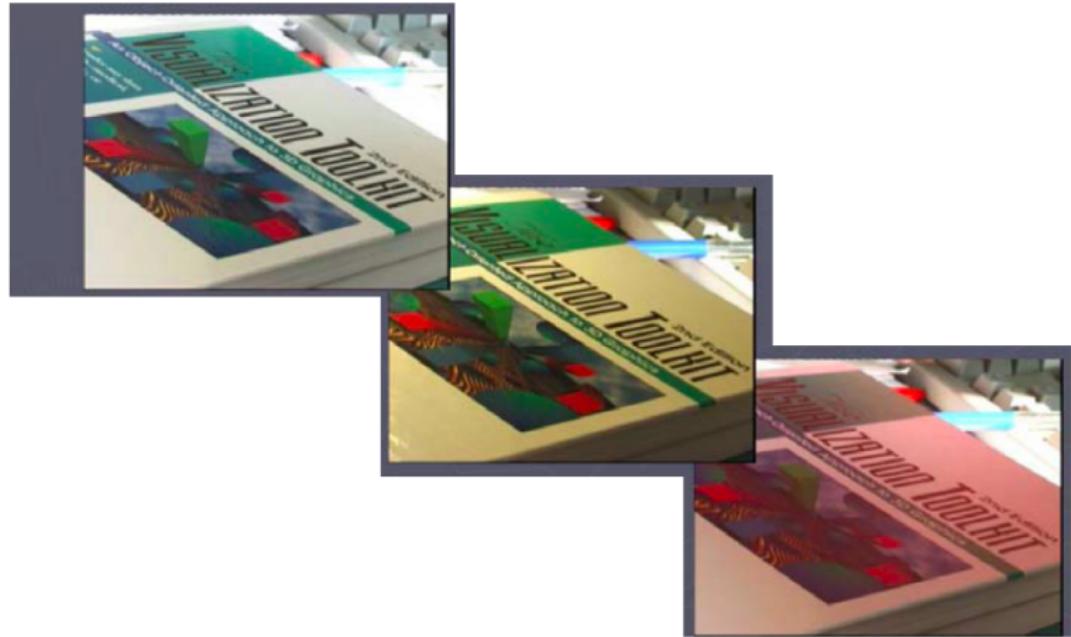
Invariances



e.g. scale,
translation,
rotation

[Source: T. Tuytelaars]

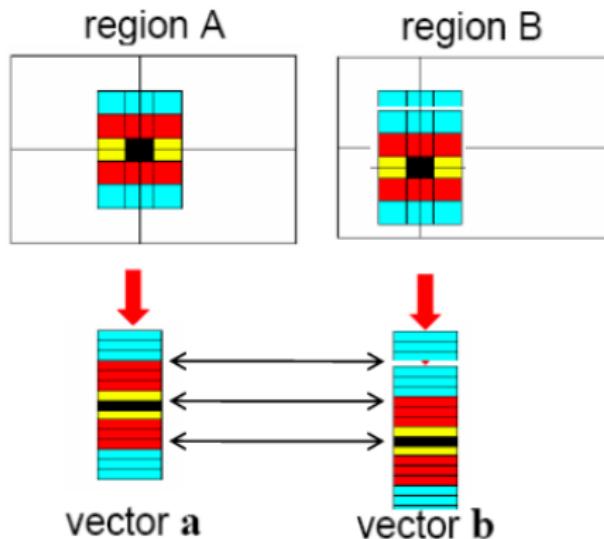
Invariances



[Source: T. Tuytelaars]

What If We Just Took Pixels?

- The simplest way is to write down the list of intensities to form a feature vector, and normalize them (i.e., mean 0, variance 1).
- Why normalization?
- But this is very sensitive to even small shifts, rotations and any affine transformation.



Tons Of Better Options

- SIFT
- PCA-SIFT
- GLOH
- HOG
- SURF
- DAISY
- LBP
- Shape Contexts
- Color Histograms

Tons Of Better Options

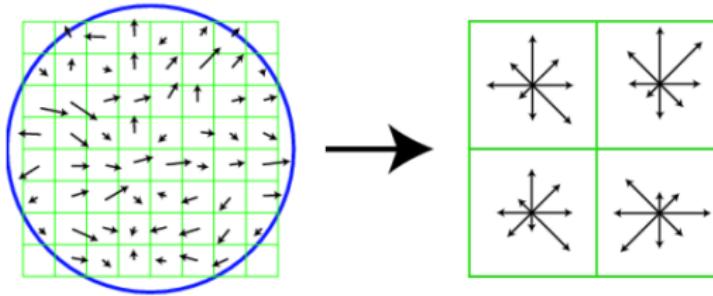
- SIFT **TODAY**
- PCA-SIFT
- GLOH
- HOG
- SURF
- DAISY
- LBP
- Shape Contexts
- Color Histograms

SIFT Descriptor [Lowe 2004]

- SIFT stands for Scale Invariant Feature Transform
- Invented by David Lowe, who also did DoG scale invariant interest points
- Actually in the same paper, which you should read:

David G. Lowe
Distinctive image features from scale-invariant keypoints
International Journal of Computer Vision, 2004

Paper: <http://www.cs.ubc.ca/~lowe/papers/ijcv04.pdf>



(a) image gradients

(b) keypoint descriptor

SIFT Descriptor

- ① Our scale invariant interest point detector gives scale ρ for each keypoint

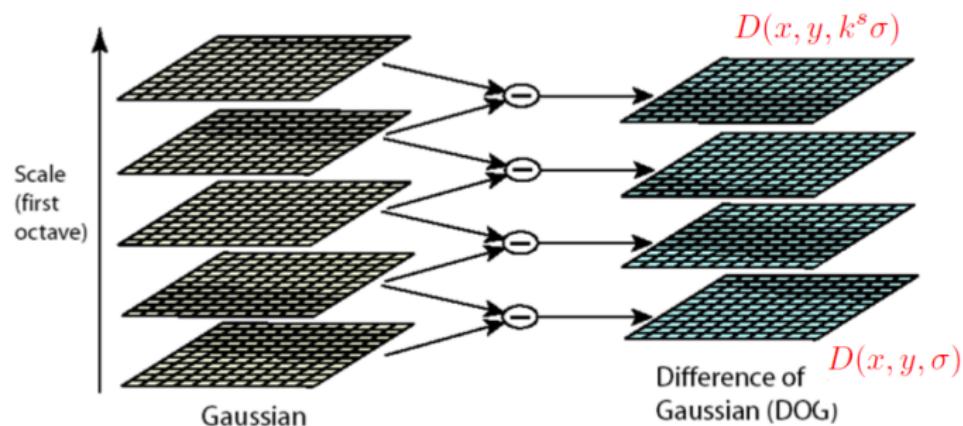
$$I_s = I * G_{k^s \sigma}$$

⋮
⋮
⋮

$$I_2 = I * G_{k^2 \sigma}$$

$$I_1 = I * G_{k\sigma}$$

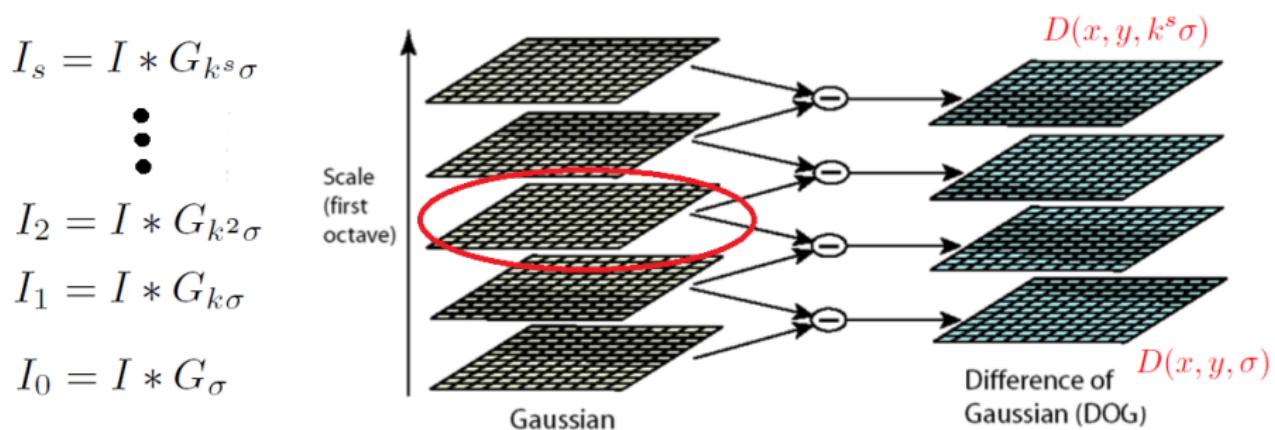
$$I_0 = I * G_\sigma$$



[Adopted from: F. Flores-Mangas]

SIFT Descriptor

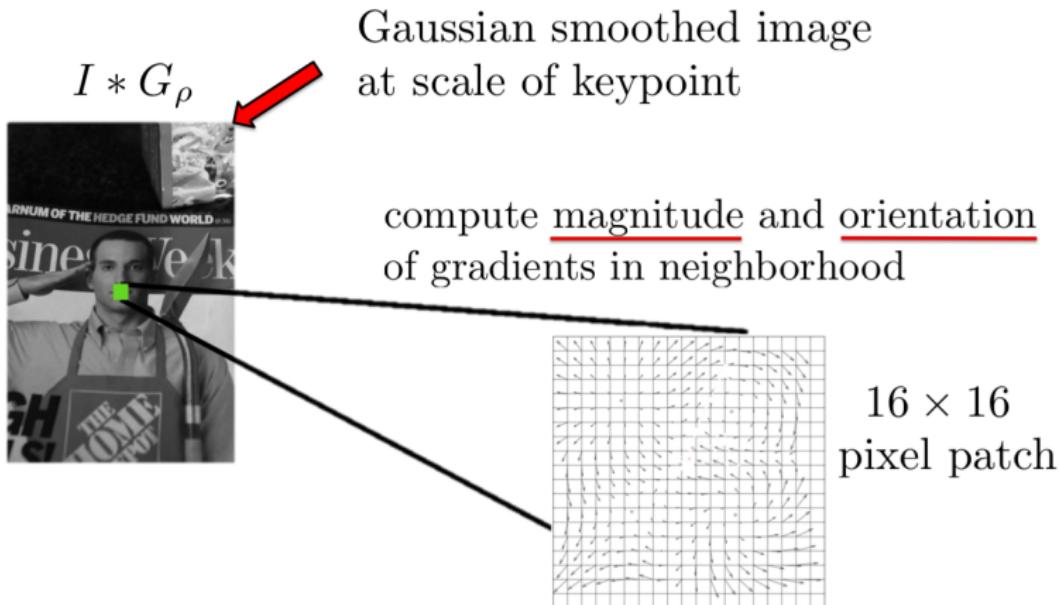
- ② For each keypoint, we take the Gaussian-blurred image at corresponding scale ρ



[Adopted from: F. Flores-Mangas]

SIFT Descriptor

- ③ Compute the gradient magnitude and orientation in neighborhood of each keypoint proportional to the detected scale



[Adopted from: F. Flores-Mangas]

SIFT Descriptor

- ③ Compute the gradient magnitude and orientation in neighborhood of each keypoint proportional to the detected scale

magnitude of gradient:

$$|\nabla I(x, y)| = \sqrt{\left(\frac{\partial(I(x, y) * G_\rho)}{\partial x}\right)^2 + \left(\frac{\partial(I(x, y) * G_\rho)}{\partial y}\right)^2}$$

gradient orientation:

$$\theta(x, y) = \arctan\left(\frac{\partial I * G_\rho}{\partial y} / \frac{\partial I * G_\rho}{\partial x}\right)$$

(in case you forgot ;))

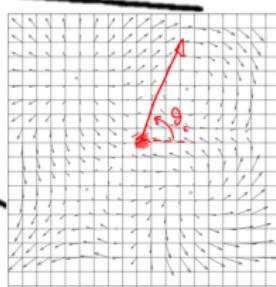
SIFT Descriptor

- ④ Compute dominant orientation of each keypoint. How?



Gaussian smoothed image
at scale of keypoint

compute magnitude and orientation
of gradients in neighborhood



16×16
pixel patch

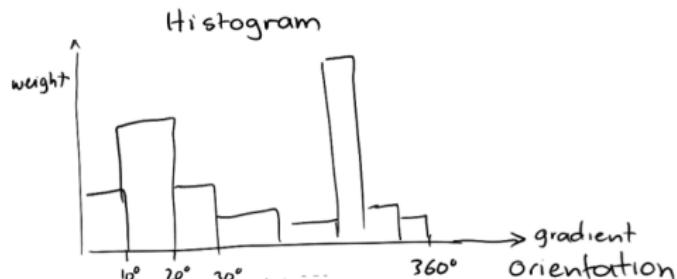
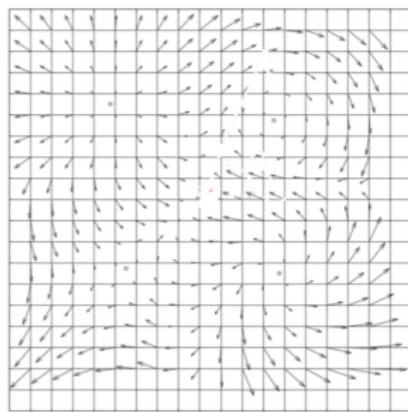
[Adopted from: F. Flores-Mangas]

SIFT Descriptor: Computing Dominant Orientation

- Compute a histogram of gradient orientations, each bin covers 10°

16 × 16

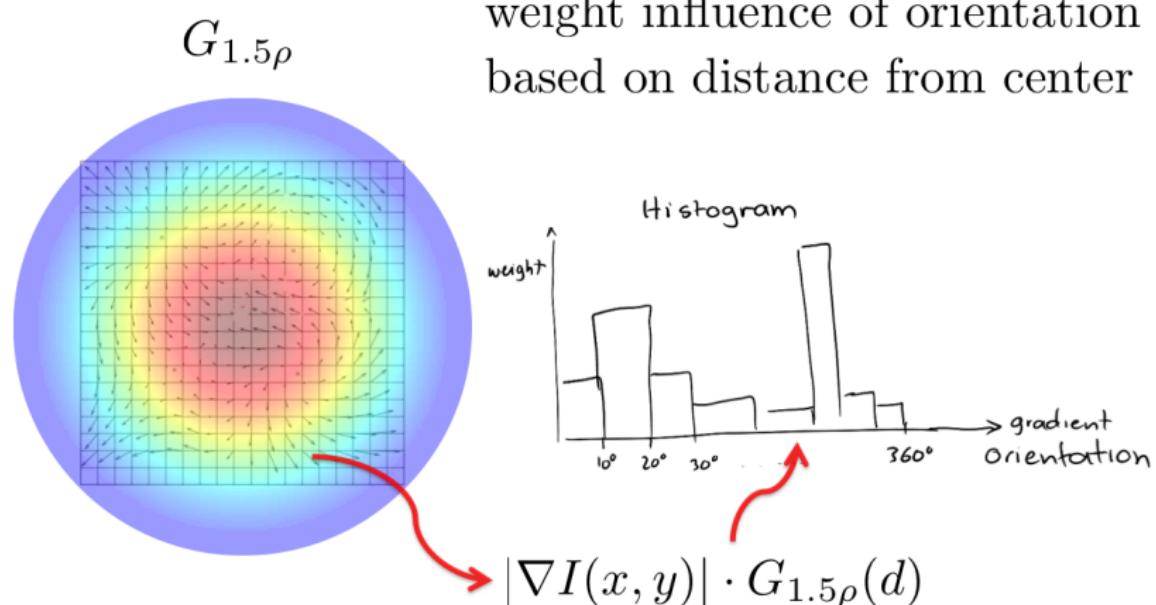
compute histograms of orientations by orientation increments of 10°



[Adopted from: F. Flores-Mangas]

SIFT Descriptor: Computing Dominant Orientation

- Compute a histogram of gradient orientations, each bin covers 10°
- Orientations closer to the keypoint center should contribute more

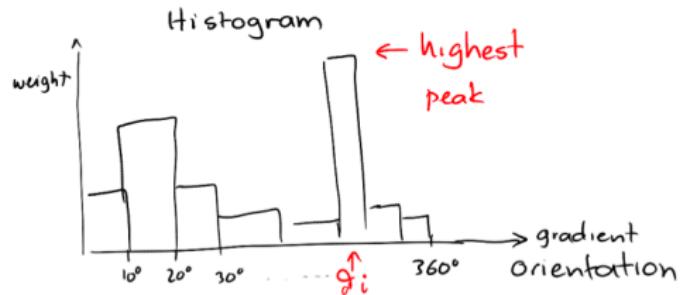
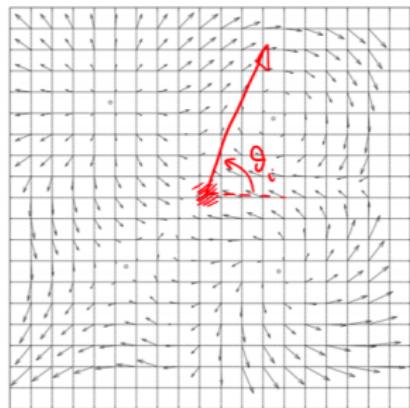


[Adopted from: F. Flores-Mangas]

SIFT Descriptor: Computing Dominant Orientation

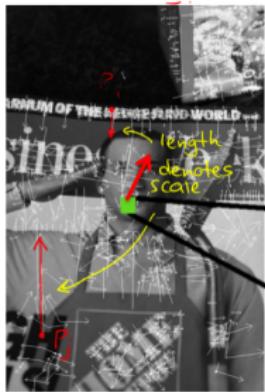
- Compute a histogram of gradient orientations, each bin covers 10°
- Orientations closer to the keypoint center should contribute more
- Orientation giving the peak in the histogram is the keypoint's orientation

16×16

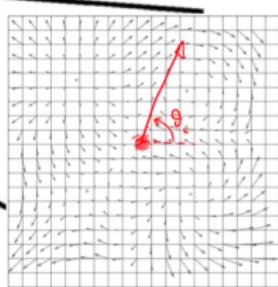


SIFT Descriptor

④ Compute dominant orientation



compute magnitude and orientation
of gradients in neighborhood



16×16
pixel patch

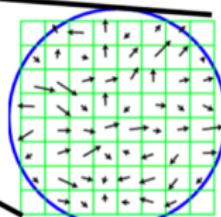
[Adopted from: F. Flores-Mangas]

SIFT Descriptor

- ⑤ Compute a 128 dimensional descriptor: 4×4 grid, each cell is a histogram of 8 orientation bins relative to dominant orientation



compute descriptor, relative
to dominant orientation



128 dim
descriptor

each descriptor has:

$$P_i = (x_i, y_i, \rho_i, \vartheta_i) \quad \text{and} \quad f_i \dots \text{128 dim vector}$$

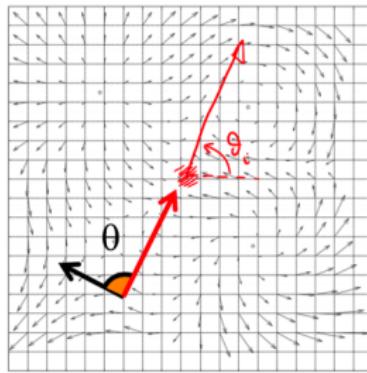
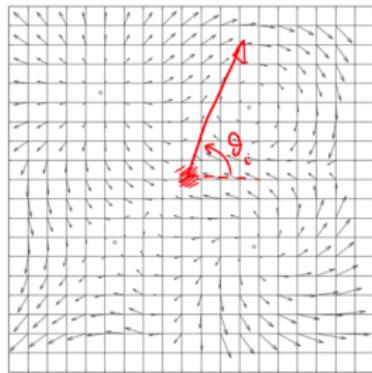
location scale orientation feature vector

[Adopted from: F. Flores-Mangas]

SIFT Descriptor: Computing the Feature Vector

- Compute the orientations **relative to the dominant orientation**
- Otherwise rotating an object would phase shift entries in histogram

16×16 patch
centered in (x_i, y_i)

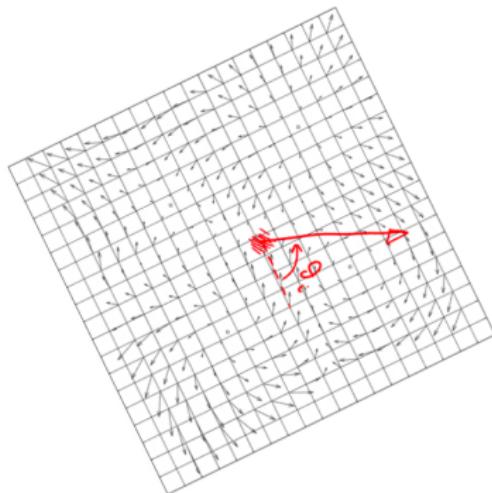
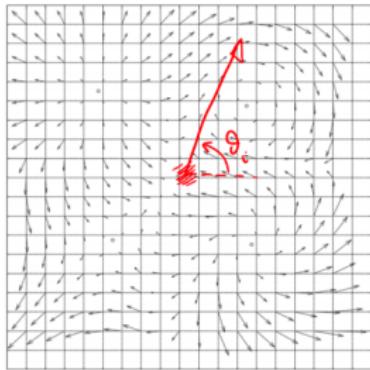


[Adopted from: F. Flores-Mangas]

SIFT Descriptor: Computing the Feature Vector

- Compute the orientations **relative to the dominant orientation**
- Otherwise rotating an object would phase shift entries in histogram

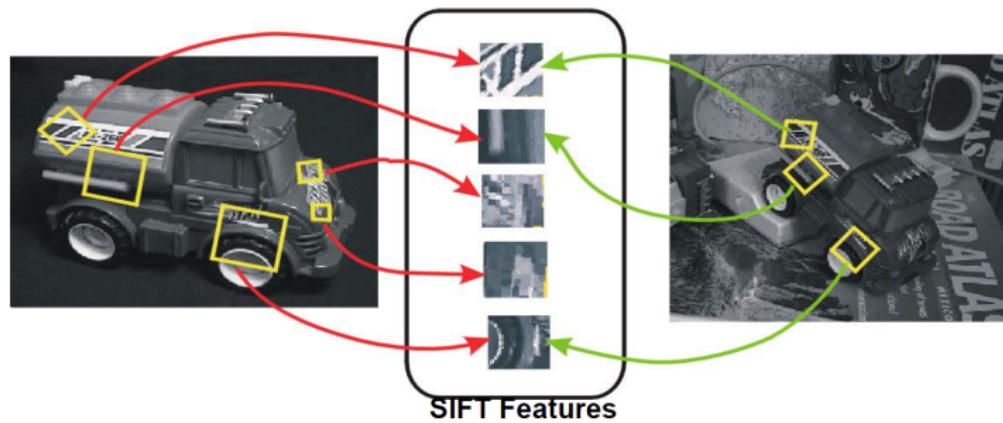
16×16 patch
centered in (x_i, y_i)



[Adopted from: F. Flores-Mangas]

SIFT Descriptor: Computing the Feature Vector

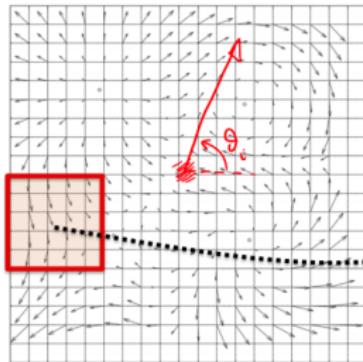
- Compute the orientations **relative** to the **dominant orientation**
- Otherwise rotating an object would phase shift entries in histogram



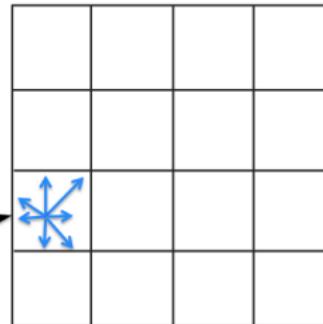
SIFT Descriptor: Computing the Feature Vector

- Compute the orientations **relative to the dominant orientation**
- Otherwise rotating an object would phase shift entries in histogram
- Form a 4×4 grid. For each grid cell compute a histogram of orientations for 8 orientation bins spaced apart by 45°

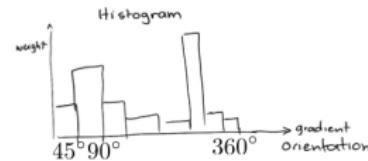
16×16 patch
centered in (x_i, y_i)



SIFT descriptor



compute histogram of orientations
this time 8 bins spaced by 45°

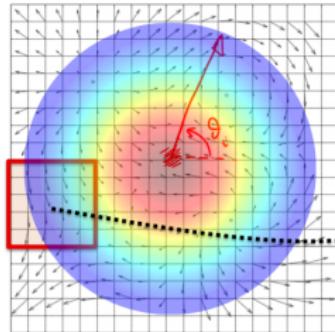


[Adopted from: F. Flores-Mangas]

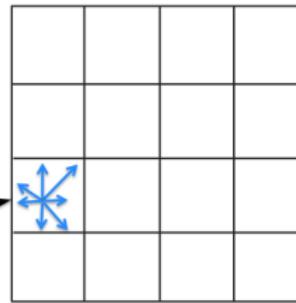
SIFT Descriptor: Computing the Feature Vector

- Compute the orientations **relative to the dominant orientation**
- Otherwise rotating an object would phase shift entries in histogram
- Form a 4×4 grid. For each grid cell compute a histogram of orientations for 8 orientation bins spaced apart by 45°

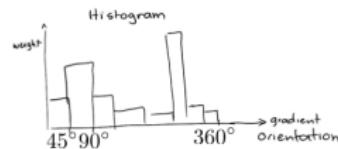
16×16 patch
centered in (x_i, y_i)



SIFT descriptor



again weigh contributions
this time: $|\nabla I(x, y)| \cdot G_{0.5\rho}$

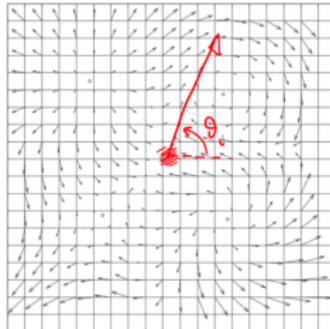


[Adopted from: F. Flores-Mangas]

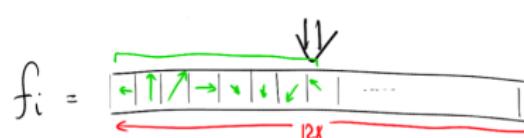
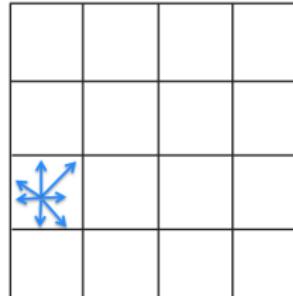
SIFT Descriptor: Computing the Feature Vector

- Compute the orientations **relative** to the **dominant orientation**
- Otherwise rotating an object would phase shift entries in histogram
- Form a 4×4 grid. For each grid cell compute a histogram of orientations for 8 orientation bins spaced apart by 45°
- Form the 128 dimensional feature vector

16×16 patch
centered in (x_i, y_i)



SIFT descriptor



[Adopted from: F. Flores-Mangas]

SIFT Descriptor: Post-processing

- The resulting 128 non-negative values form a **raw version** of the SIFT descriptor vector.
- To reduce the **effects of contrast or gain** (additive variations are already removed by the gradient), the 128-D vector is normalized to unit length: $f_i = f_i / \|f_i\|$

SIFT Descriptor: Post-processing

- The resulting 128 non-negative values form a **raw version** of the SIFT descriptor vector.
- To reduce the **effects of contrast or gain** (additive variations are already removed by the gradient), the 128-D vector is normalized to unit length: $f_i = f_i / \|f_i\|$
- To further make the descriptor robust to other **photometric variations**, values are clipped to 0.2 and the resulting vector is once again renormalized to unit length.

SIFT Descriptor: Post-processing

- The resulting 128 non-negative values form a **raw version** of the SIFT descriptor vector.
- To reduce the **effects of contrast or gain** (additive variations are already removed by the gradient), the 128-D vector is normalized to unit length: $f_i = f_i / \|f_i\|$
- To further make the descriptor robust to other **photometric variations**, values are clipped to 0.2 and the resulting vector is once again renormalized to unit length.
- Great engineering effort!

SIFT Descriptor: Post-processing

- The resulting 128 non-negative values form a **raw version** of the SIFT descriptor vector.
- To reduce the **effects of contrast or gain** (additive variations are already removed by the gradient), the 128-D vector is normalized to unit length: $f_i = f_i / \|f_i\|$
- To further make the descriptor robust to other **photometric variations**, values are clipped to 0.2 and the resulting vector is once again renormalized to unit length.
- Great engineering effort!
- What is SIFT invariant to?

SIFT Descriptor: Post-processing

- The resulting 128 non-negative values form a **raw version** of the SIFT descriptor vector.
- To reduce the **effects of contrast or gain** (additive variations are already removed by the gradient), the 128-D vector is normalized to unit length: $f_i = f_i / \|f_i\|$
- To further make the descriptor robust to other **photometric variations**, values are clipped to 0.2 and the resulting vector is once again renormalized to unit length.
- Great engineering effort!
- What is SIFT invariant to?

Properties of SIFT

Invariant to:

- Scale
- Rotation

Partially invariant to:

- Illumination changes (sometimes even day vs. night)
- Camera viewpoint (up to about 60 degrees of out-of-plane rotation)
- Occlusion, clutter (why?)

Also important:

- Fast and efficient – can run in real time
- Lots of code available

Examples



Figure: Matching in day / night under viewpoint change

[Source: S. Seitz]

Example

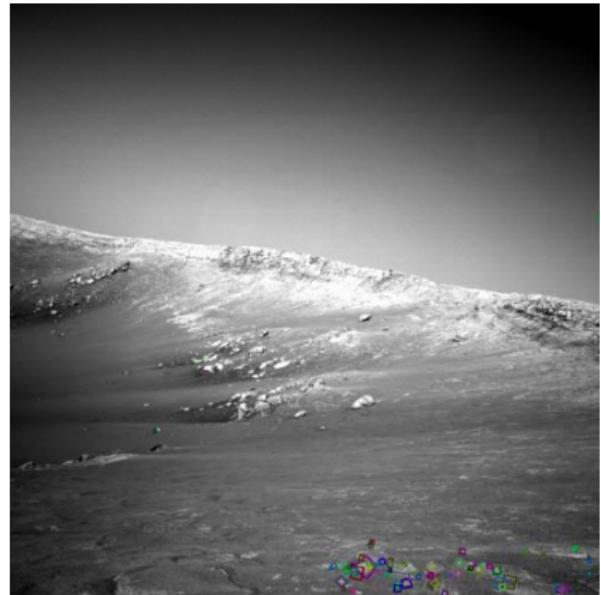


Figure: NASA Mars Rover images with SIFT feature matches

[Source: N. Snavely]

PCA-SIFT

- The dimensionality of SIFT is pretty high, i.e., 128D for each keypoint
- Reduce the dimensionality using linear dimensionality reduction
- In this case, principal component analysis (PCA)
- Use 10D or so descriptor

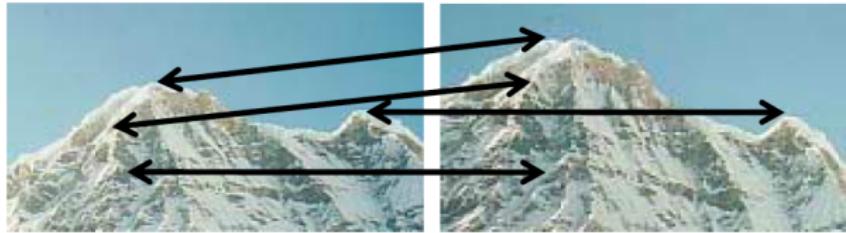
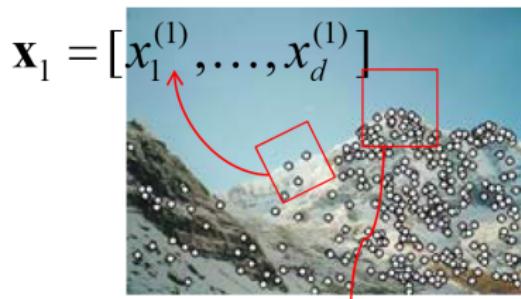
[Source: R. Urtasun]

Other Descriptors

- SURF
- DAISY
- LBP
- HOG
- Shape Contexts
- Color Histograms

Local Features

- **Detection:** Identify the interest points.
- **Description:** Extract feature descriptor around each interest point.
- **Matching:** Determine correspondence between descriptors in two views.



[Source: K. Grauman]

Image Features: Matching the Local Descriptors

Matching the Local Descriptors

Once we have extracted keypoints and their descriptors, we want to match the features between pairs of images.

- Ideally a match is a correspondence between a local part of the object on one image to the same local part of the object in another image
- How should we compute a match?



Figure: Images from K. Grauman

Matching the Local Descriptors

Once we have extracted keypoints and their descriptors, we want to match the features between pairs of images.

- Ideally a match is a correspondence between a local part of the object on one image to the same local part of the object in another image
- How should we compute a match?

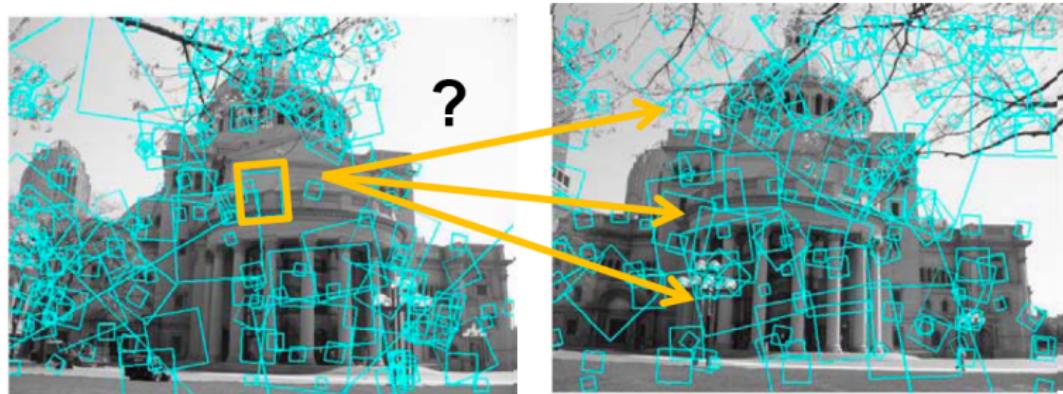
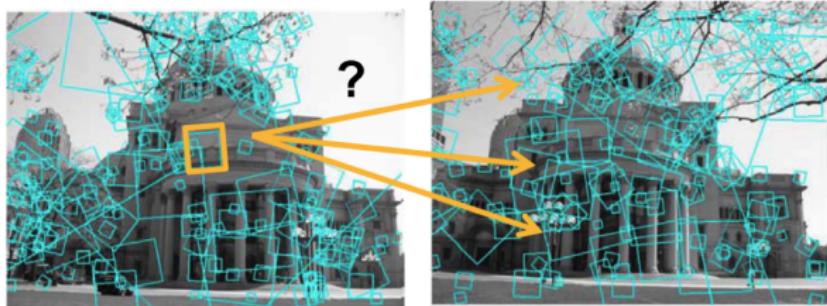


Figure: Images from K. Grauman

Matching the Local Descriptors

- Simple: **Compare them all**, compute Euclidean distance



$$f_1 \quad \boxed{\text{---}}$$

$$f_2 \quad \boxed{\text{---}}$$

$$f_3 \quad \boxed{\text{---}}$$

$$f'_1 \quad \boxed{\text{---}}$$

$$f'_2 \quad \boxed{\text{---}}$$

$$f'_3 \quad \boxed{\text{---}}$$

$$f_{k-1} \quad \boxed{\text{---}}$$

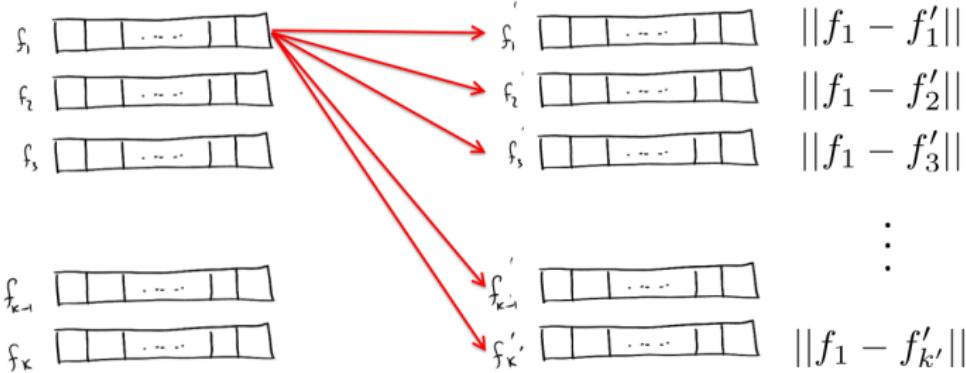
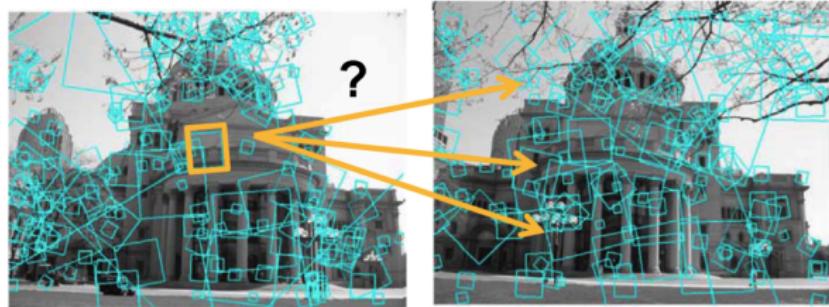
$$f_k \quad \boxed{\text{---}}$$

$$f'_{k-1} \quad \boxed{\text{---}}$$

$$f'_{k'} \quad \boxed{\text{---}}$$

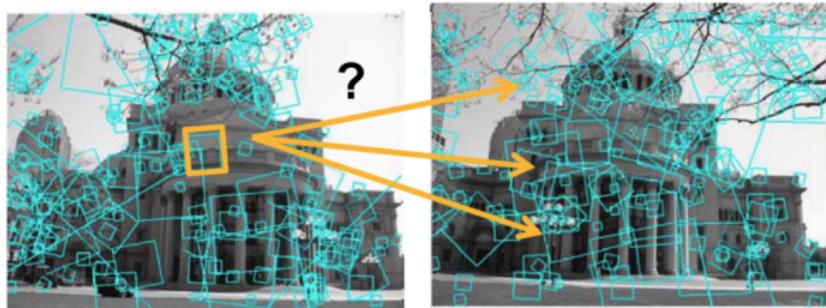
Matching the Local Descriptors

- Simple: **Compare them all**, compute Euclidean distance



Matching the Local Descriptors

- Find closest match (min distance). How do we know if match is **reliable**?



$$f_1 \quad \boxed{\begin{array}{|c|c|c|c|} \hline & & & \\ \hline \end{array}}$$
$$f_2 \quad \boxed{\begin{array}{|c|c|c|c|} \hline & & & \\ \hline \end{array}}$$
$$f_3 \quad \boxed{\begin{array}{|c|c|c|c|} \hline & & & \\ \hline \end{array}}$$

$$f_{k-1} \quad \boxed{\begin{array}{|c|c|c|c|} \hline & & & \\ \hline \end{array}}$$
$$f_k \quad \boxed{\begin{array}{|c|c|c|c|} \hline & & & \\ \hline \end{array}}$$

$$f'_1 \quad \boxed{\begin{array}{|c|c|c|c|} \hline & & & \\ \hline \end{array}} \quad \|f_1 - f'_1\|$$

$$f'_2 \quad \boxed{\begin{array}{|c|c|c|c|} \hline & & & \\ \hline \end{array}} \quad \|f_1 - f'_2\|$$

$$f'_3 \quad \boxed{\begin{array}{|c|c|c|c|} \hline & & & \\ \hline \end{array}} \quad \|f_1 - f'_3\|$$

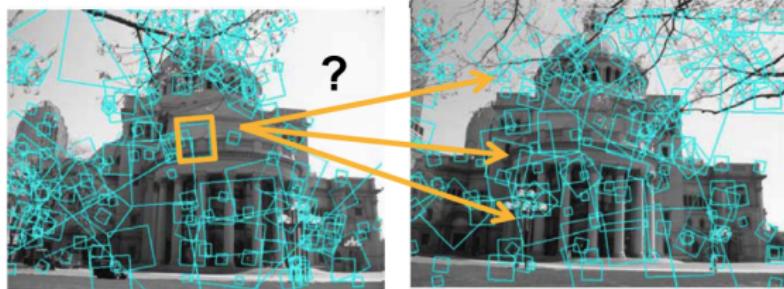
min (closest match)

$$f'_{k-1} \quad \boxed{\begin{array}{|c|c|c|c|} \hline & & & \\ \hline \end{array}}$$

$$f'_k \quad \boxed{\begin{array}{|c|c|c|c|} \hline & & & \\ \hline \end{array}} \quad \|f_1 - f'_k\|$$

Matching the Local Descriptors

- Find also the second closest match. Match reliable if first distance “much” smaller than second distance



$$f_1 \quad \boxed{\begin{array}{|c|c|c|c|} \hline & & & \\ \hline & & & \\ \hline & \dots & \dots & \\ \hline \end{array}}$$

$$f_2 \quad \boxed{\begin{array}{|c|c|c|c|} \hline & & & \\ \hline & & & \\ \hline & \dots & \dots & \\ \hline \end{array}}$$

$$f_3 \quad \boxed{\begin{array}{|c|c|c|c|} \hline & & & \\ \hline & & & \\ \hline & \dots & \dots & \\ \hline \end{array}}$$

$$f_{k-1} \quad \boxed{\begin{array}{|c|c|c|c|} \hline & & & \\ \hline & & & \\ \hline & \dots & \dots & \\ \hline \end{array}}$$

$$f_k \quad \boxed{\begin{array}{|c|c|c|c|} \hline & & & \\ \hline & & & \\ \hline & \dots & \dots & \\ \hline \end{array}}$$

$$f'_1 \quad \boxed{\begin{array}{|c|c|c|c|} \hline & & & \\ \hline & & & \\ \hline & \dots & \dots & \\ \hline \end{array}} \quad \|f_1 - f'_1\|$$

$$f'_2 \quad \boxed{\begin{array}{|c|c|c|c|} \hline & & & \\ \hline & & & \\ \hline & \dots & \dots & \\ \hline \end{array}} \quad \|f_1 - f'_2\|$$

$$f'_3 \quad \boxed{\begin{array}{|c|c|c|c|} \hline & & & \\ \hline & & & \\ \hline & \dots & \dots & \\ \hline \end{array}} \quad \|f_1 - f'_3\|$$

min (closest match)

$$f'_{k-1} \quad \boxed{\begin{array}{|c|c|c|c|} \hline & & & \\ \hline & & & \\ \hline & \dots & \dots & \\ \hline \end{array}} \quad \|f_1 - f'_{k-1}\|$$

$$f'_k \quad \boxed{\begin{array}{|c|c|c|c|} \hline & & & \\ \hline & & & \\ \hline & \dots & \dots & \\ \hline \end{array}} \quad \|f_1 - f'_k\|$$

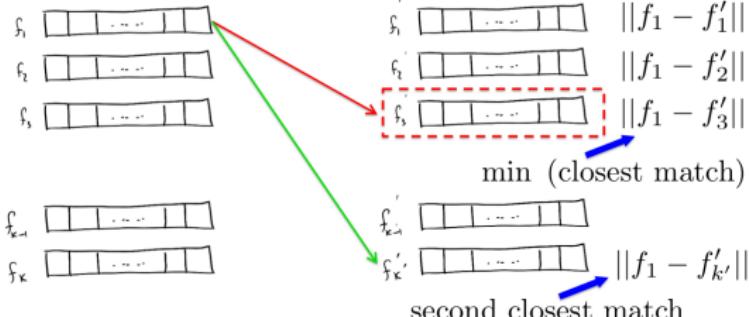
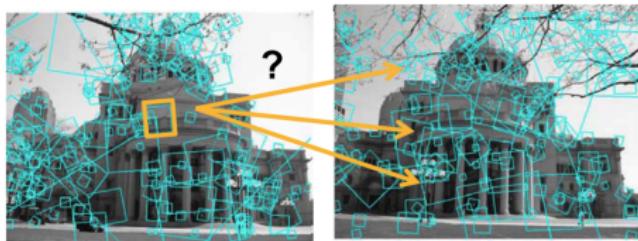
second closest match

Matching the Local Descriptors

- Compute the ratio:

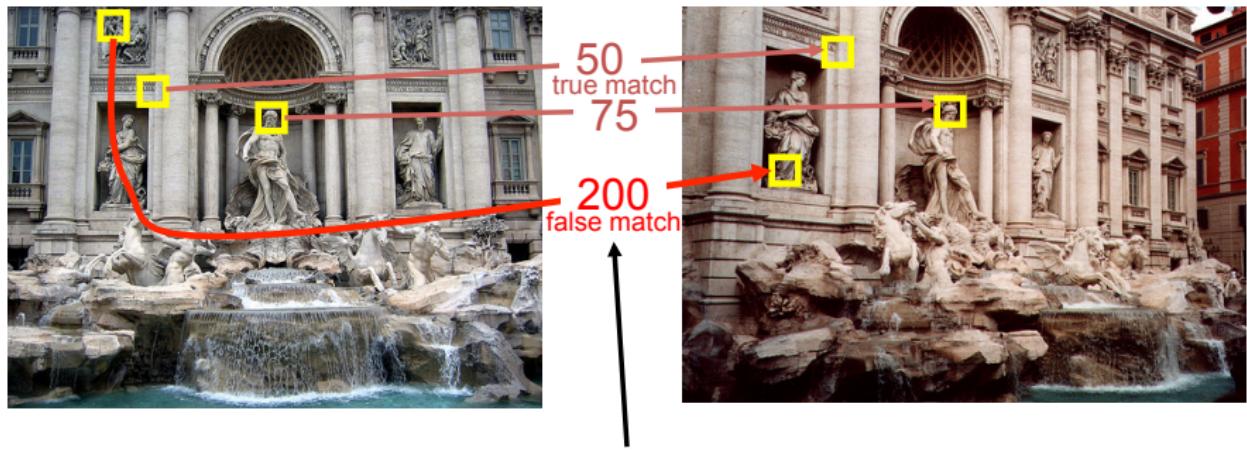
$$\phi_i = \frac{\|f_i - f'_i^*\|}{\|f_i - f'^{**}_i\|}$$

where f'_i^* is the closest and f'^{**}_i second closest match to f_i .



Which Threshold to Use?

- Setting the threshold too high results in too many false positives, i.e., incorrect matches being returned.
- Setting the threshold too low results in too many false negatives, i.e., too many correct matches being missed



[Source: N. Snavely]

Babak Taati

CSC420: Intro to Image Understanding

23 / 30

Which Threshold to Use?

- Threshold ratio of nearest to 2nd nearest descriptor
- Typically: $\phi_i < 0.8$

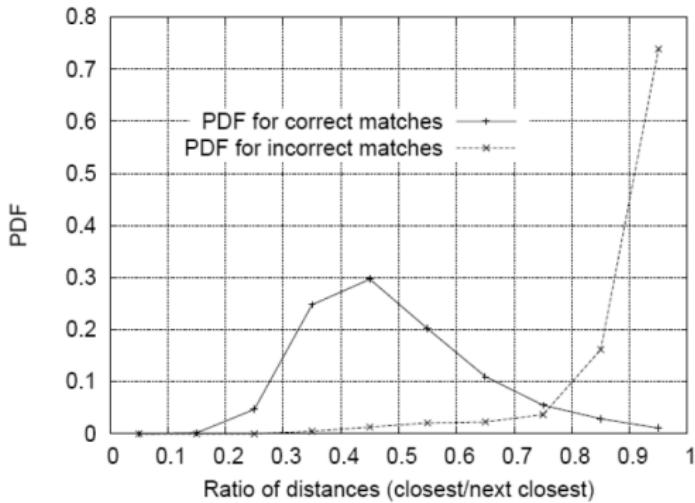


Figure: Images from D. Lowe

[Source: K. Grauman]

Applications of Local Invariant Features

- Wide baseline stereo
- Motion tracking
- Panorama stitching
- Mobile robot navigation
- 3D reconstruction
- Recognition
- Retrieval

[Source: K. Grauman]

Wide Baseline Stereo



[Source: T. Tuytelaars]

Motion Tracking



Figure: Images from J. Pilet

Now What

- Now we know how to extract scale and rotation invariant features
- We even know how to match features across images
- Can we use this to find Waldo in an even more sneaky scenario?

Now What

- Now we know how to extract scale and rotation invariant features
- We even know how to match features across images
- Can we use this to find Waldo in an even more sneaky scenario?



Waldo on the road



template

Now What

- Now we know how to extract scale and rotation invariant features
- We even know how to match features across images
- Can we use this to find Waldo in an even more sneaky scenario?



template

He comes closer... We know how to solve this

Now What

- Now we know how to extract scale and rotation invariant features
- We even know how to match features across images
- Can we use this to find Waldo in an even more sneaky scenario?

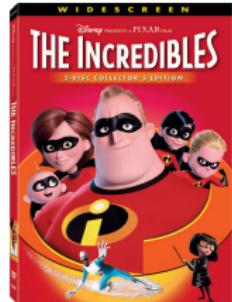


template

Someone takes a (weird) picture of him!

Find My DVD!

- More interesting: If we have DVD covers (e.g., from Amazon), can we match them to DVDs in real scenes?



Matching Planar Objects In New Viewpoints