

Object Detection

But first,
Course Evaluation (on Quercus)

Type of Approaches

Different approaches tackle detection differently. They can roughly be categorized into three main types:

- Find **interest points**, followed by Hough voting
- **Sliding windows**: “slide” a box around image and classify each image crop inside a box (contains object or not?)
- Generate **region (object) proposals**, and classify each region ← We have looked at R-CNN a little bit. Today we'll focus on how to group pixels in super pixels, and regions. Once you have a region, you just attack it with a Neural Network.

Segmentation

Segmentation

- Let's start with a very general definition of segmentation as an optimization problem
- Define an image, $\mathbf{I} = I(x, y)$ for $x, y \in \Omega$
- A segmentation is a labelling of every pixel into an object class, $\mathbf{S} = S(x, y)$ for $x, y \in \Omega$
- As you saw in assignment 4, to segment the car you first had to know where it is. But segmenting it could also give you a better estimate of its centre.

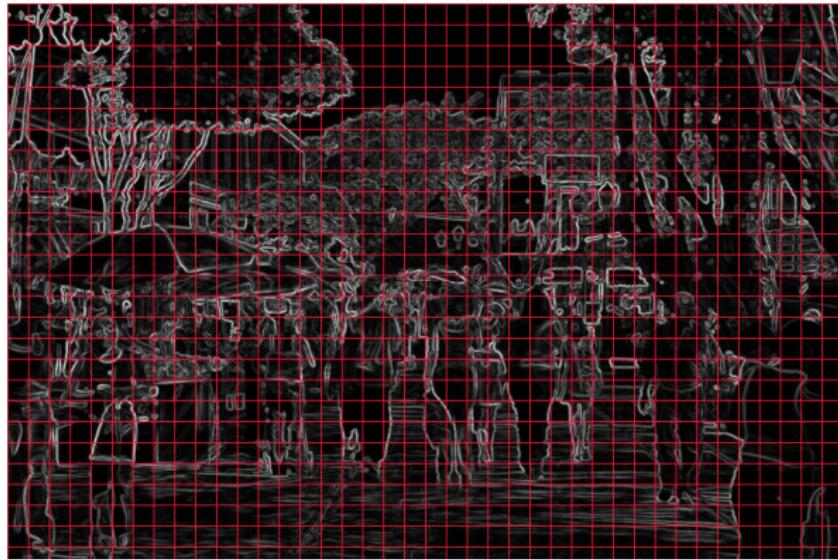
How Can We Get Semantics Quickly?

- Each image has about 1 to 4 **million pixels**.
- That's a lot.
- A real robotics system needs to recognize things really fast if it wants to react to the world around it in real-time
- And secondly, pixels might not be the best level of organization to capture meaningful semantics of the object



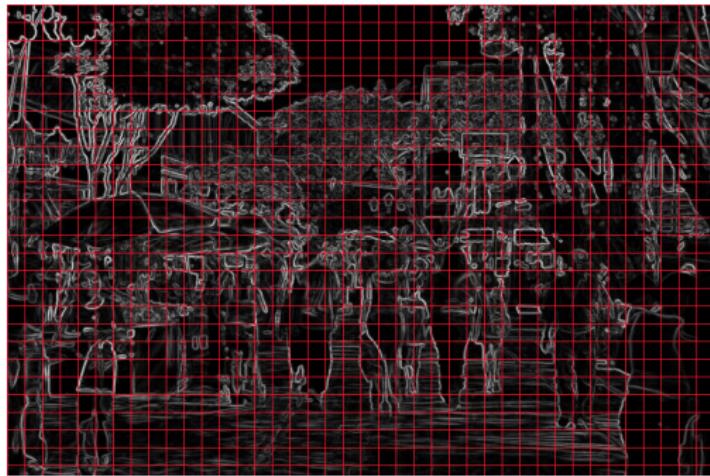
Remember HOG?

- HOG divides an image into cells. There is 64-times less cells than pixels. Two orders of magnitude less. Yet still, there is a lot of them.

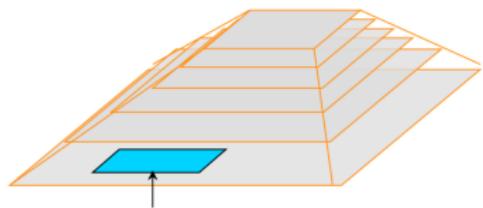


Remember HOG?

- HOG divides an image into cells. There is 64-times less cells than pixels. Two orders of magnitude less. Yet still, there is a lot of them.
- Let's not forget we still need to run the detector, not only tonnes of locations, but also tonnes of scales.
- And who said rectangular grid was the best choice. It was only convenient!



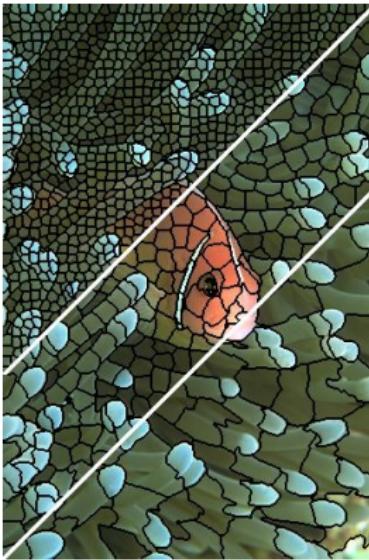
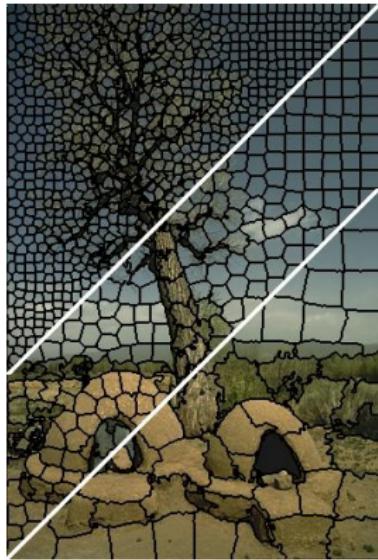
Scale-space pyramid



Detection window

Superpixels

- One of the ideas is to merge similar pixels into (way fewer) “superpixels”.
- **First question** to ask: How does that help us?
- **Second question:** What properties should superpixels (or any regions) satisfy in order to be useful?



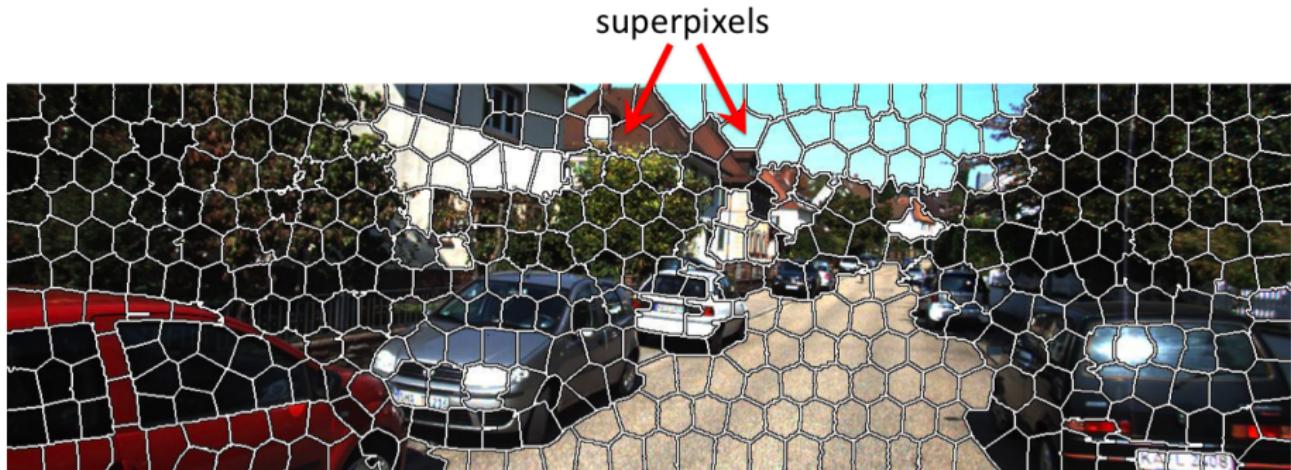
Why Are Superpixels Useful?

- Example 1: How can we find all *road* pixels in this image?



Why Are Superpixels Useful?

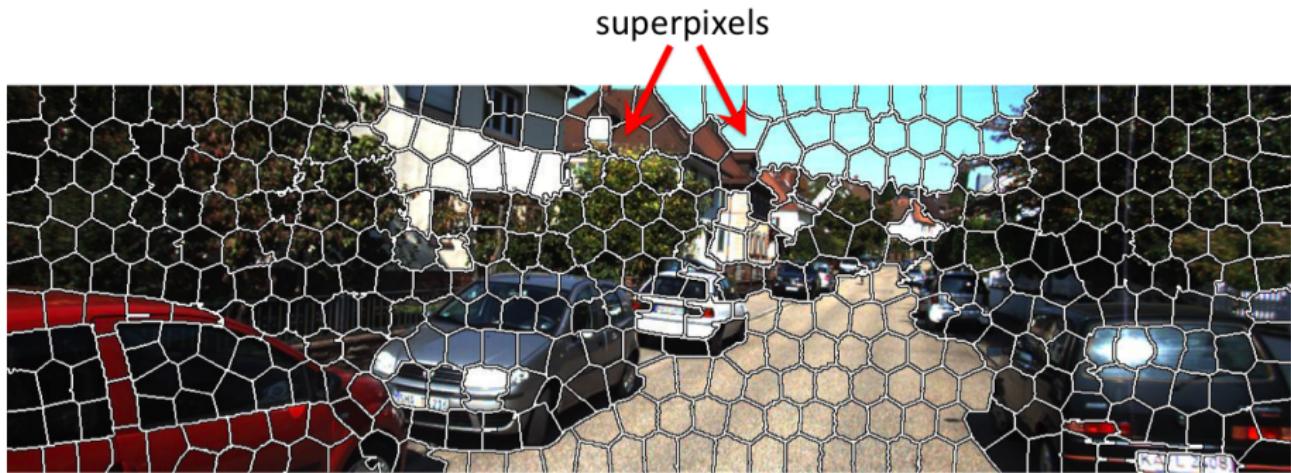
- Compute superpixels. A 4million pixel image is converted into only 500 superpixels. And now?



[Superpixels computed by Jian Yao, PhD student at UofT. Thanks Jian!]

Why Are Superpixels Useful?

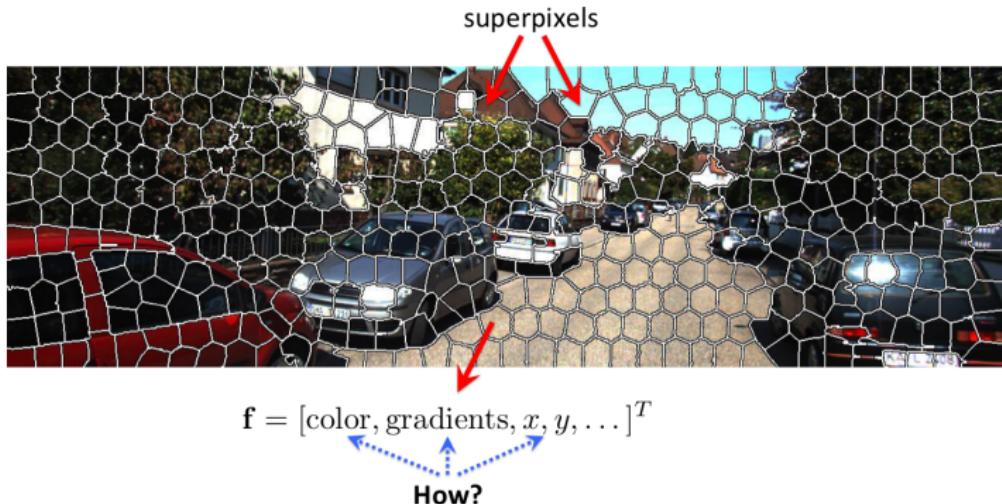
- Possible idea: Compute features on each superpixel and train a classifier for road/non-road. Use this classifier at test time



[Superpixels computed by Jian Yao, PhD student at UofT. Thanks Jian!]

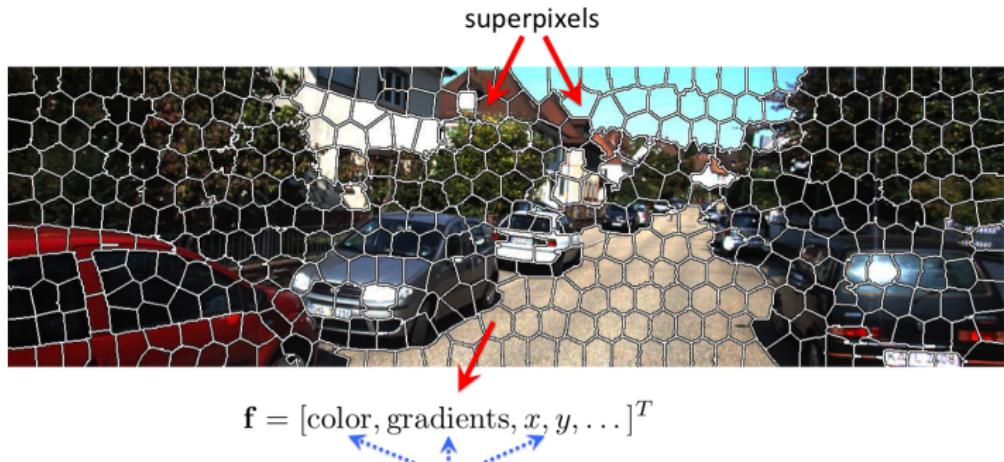
Why Are Superpixels Useful?

- Compute features for each superpixel. (make sure to normalize them, e.g, to norm or max value 1; classifiers will work better)
- Different superpixels have different number of pixels. How can I compute my feature vector (has to have the same dimension for each superpixel)?



Why Are Superpixels Useful?

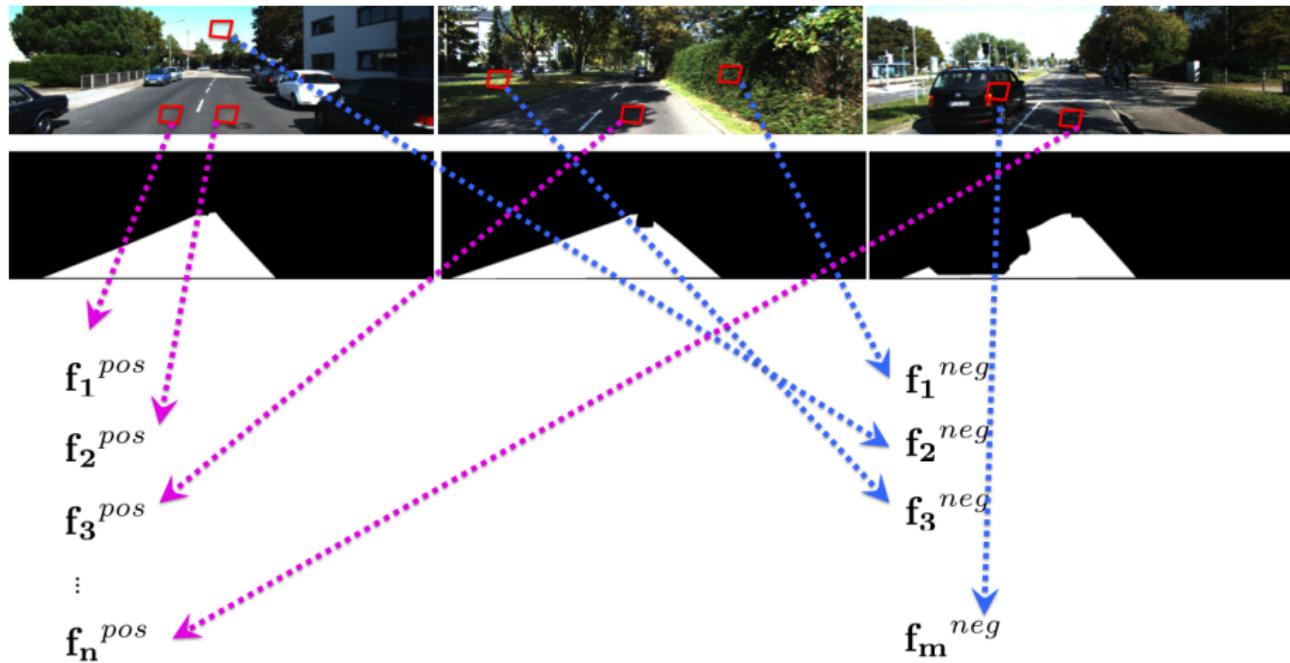
- Average or better: **histogram** (typically the more dimensions that the feature vector has, the better the classifiers work). Histograms allow you to inflate a feature to multiple dimensions. And now?



Can be: 1) average across pixels in superpixel, 2) a histogram, 3) a histogram of visual words

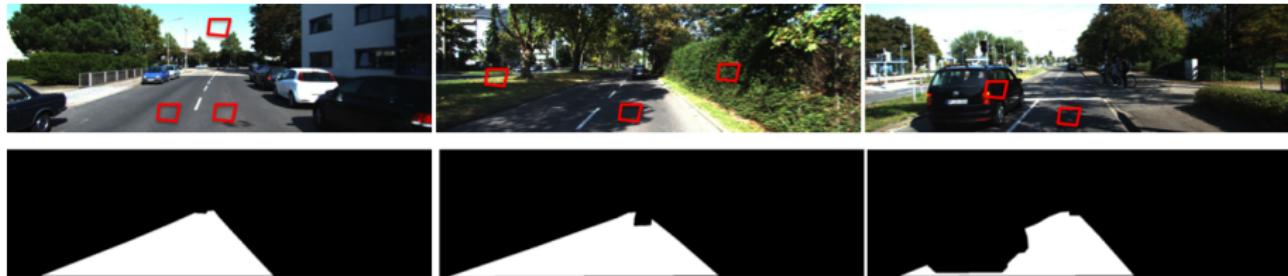
Why Are Superpixels Useful?

- Collect (randomly sample) a bunch of positive features (from regions annotated as *road*) and negative features (from regions of non-road).



Why Are Superpixels Useful?

- Train your favorite classifier. And at test time?



\mathbf{f}_1^{pos}

\mathbf{f}_2^{pos}

\mathbf{f}_3^{pos}

:

\mathbf{f}_n^{pos}

Train classifier

$$\mathbf{w}^T \cdot \mathbf{f} + b = 0$$

\mathbf{f}_1^{neg}

\mathbf{f}_2^{neg}

\mathbf{f}_3^{neg}

\mathbf{f}_m^{neg}

Why Are Superpixels Useful?

- Some problems can emerge, e.g., shadows, or grayish things on buildings that can confuse the classifier



Shadows look like anything else (car, dark tree, etc). Prediction will be bad.
Can we do something about it?

Why Are Superpixels Useful?

- We can make use of **depth** (e.g, road is typically not 1m above our eye level). Which brings us to:
- Example 2: We can use superpixels also for de-noising stereo results

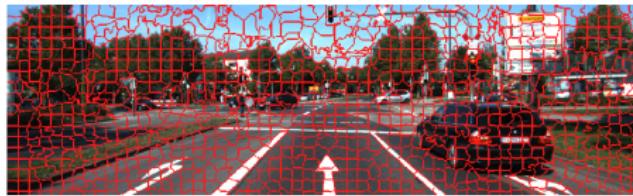
(a) left image



(b) right image



(c) superpixels



(d) estimated disparity

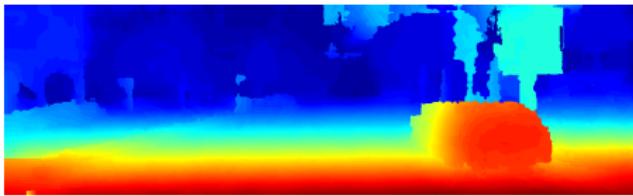


Figure: <http://ttic.uchicago.edu/~dmcalester/SPS/index.html>

Why Are Superpixels Useful?

- (back to) Example 1: Now just simply augment superpixel features with 3D features.
- And hope for the best.



$$\mathbf{f} = [\text{color, gradients, } x, y, Y, \text{3d features (gradients in 3D?), ...}]$$



Add depth informed features

Why Are Superpixels Useful?

- RGB-D images

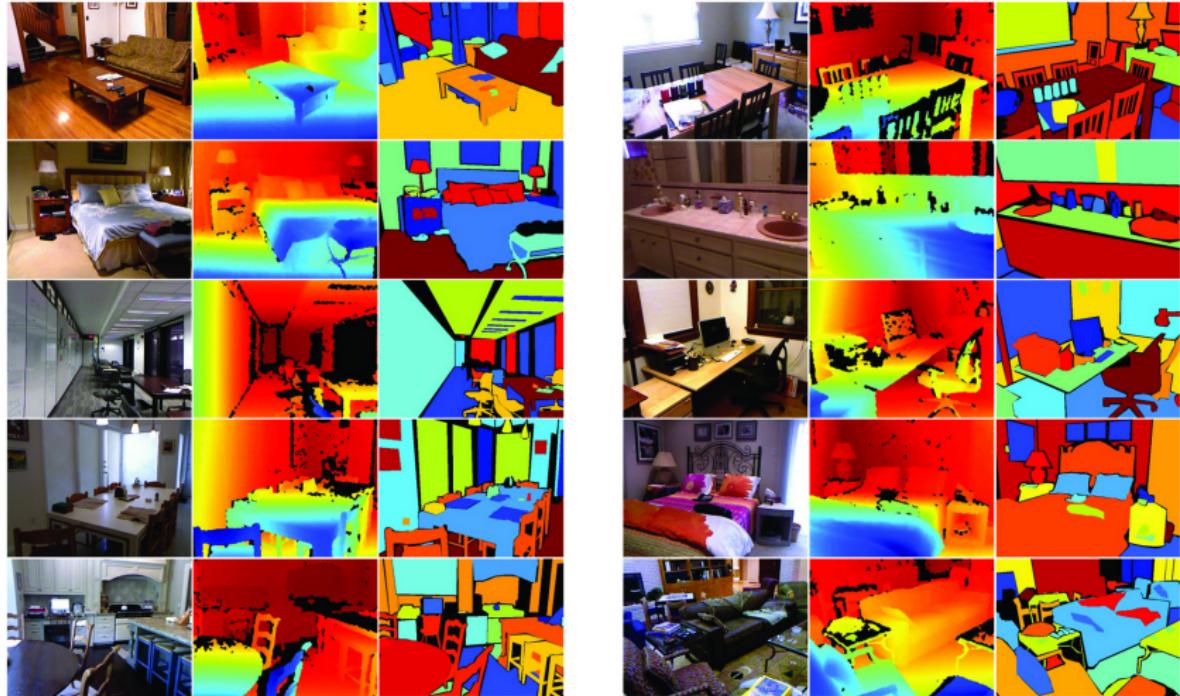
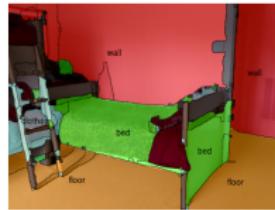
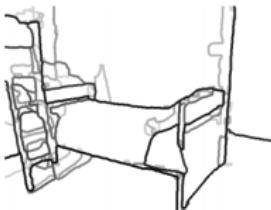
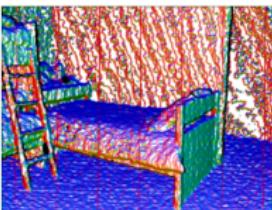
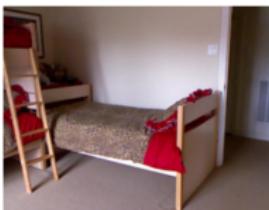
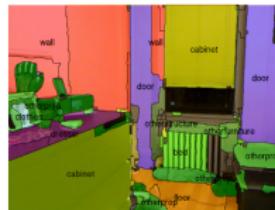
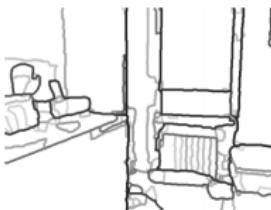
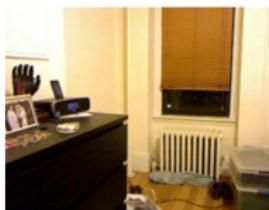
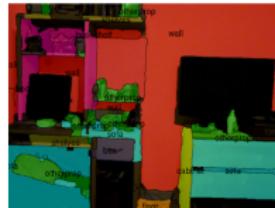


Figure: http://cs.nyu.edu/~silberman/datasets/nyu_depth_v2.html

Why Are Superpixels Useful?

- Example 2: You can also train superpixel classifiers with RGB-D features to predict a variety of **semantic classes**



(a) Color Image

(b) Depth

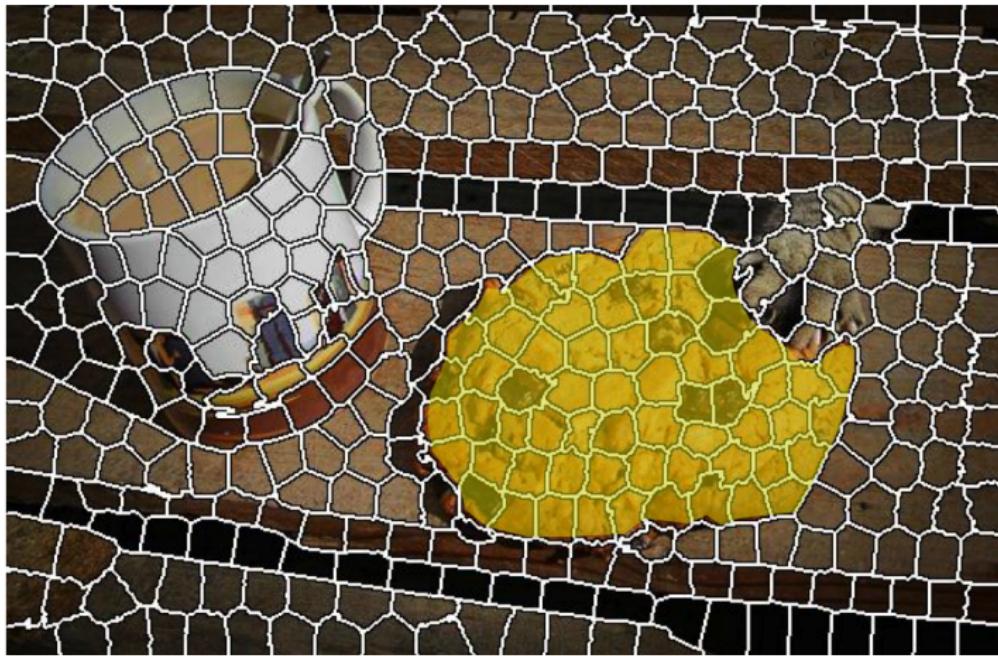
(c) Contours

(d) Semantic Segmentation

Figure: <http://www.cs.berkeley.edu/~sgupta/pdf/GuptaArbelaezMalikCVPR13.pdf>

Why Are Superpixels Useful?

- Example 3: And of course also to find our cookie.



[Superpixels computed by Jian Yao, PhD student at UofT. Thanks Jian!]

Why Are Superpixels Useful?

- Example 7: We can use them for tracking. And many more things.



Figure:

http://groups.csail.mit.edu/vision/sli/projects.php?name=temporal_superpixels

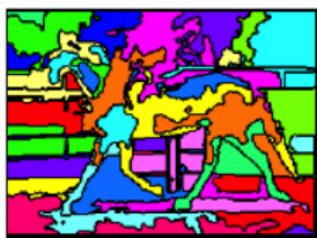
Segmentation

- Question 2: What properties should our segmentation have?

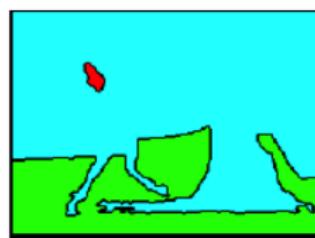
Segmentation

- Should be fast to compute
- Should not merge different objects (undersegmentation)

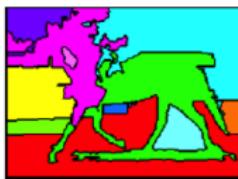
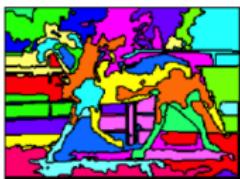
Types of segmentations



Oversegmentation



Undersegmentation



Multiple Segmentations

Slide by D. Hoiem

Over Segmentation Algorithms

There are a lot of them out there

- Felzenswalb and Huttenlocher's Graph-based Segmentation → Let's quickly look at this one
- SLIC
- SEEDS
- Shi and Malik's Normalized Cuts
- Malik's group: Probability of boundary (gPb)
- Grundmann et al: Hierarchical Graph-based Video Segmentation
- Chang et al., Temporal Superpixels

For most the code is available!

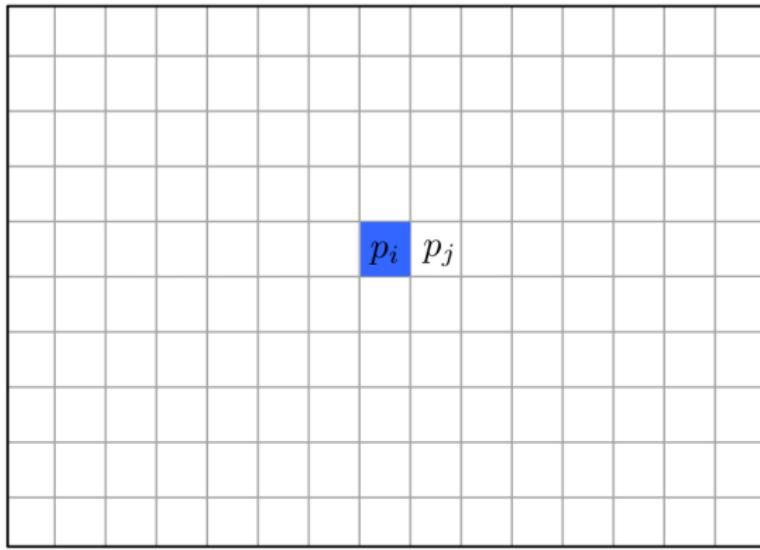
Felzenswalb and Huttenlocher's Graph-based Segmentation:

<http://cs.brown.edu/~pff/segment/>

Felzenswalb & Huttenlocher

- Take an image.

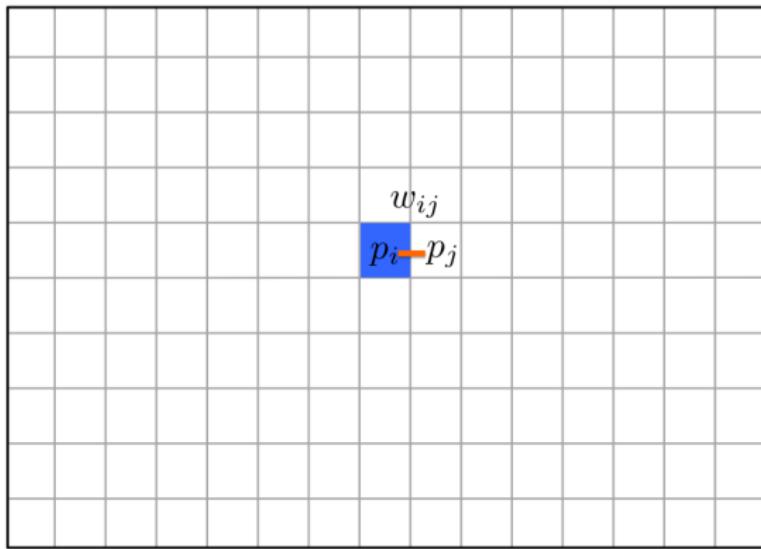
I



Felzenswalb & Huttenlocher

- Compute a weight between a pair of neighboring pixels. A weight reflects **dissimilarity**.

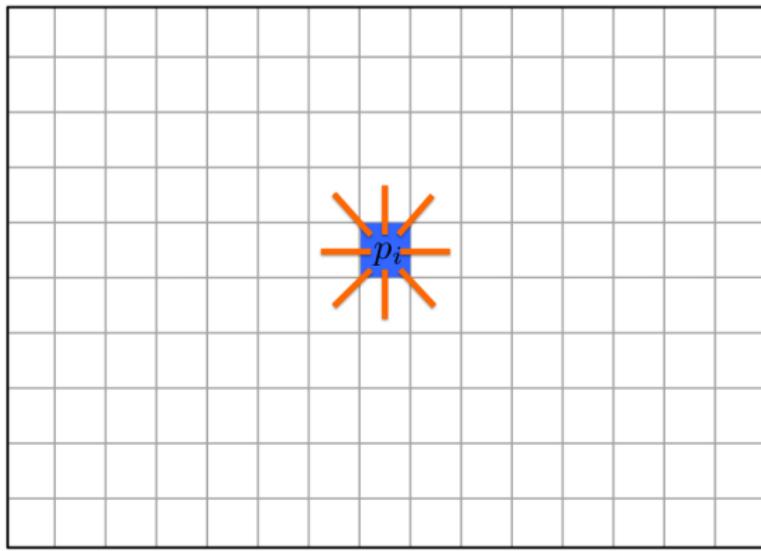
$$I \qquad w_{ij} = |I(p_i) - I(p_j)|$$



Felzenswalb & Huttenlocher

- Compute weights between a pixel and all 8 of its neighbors.

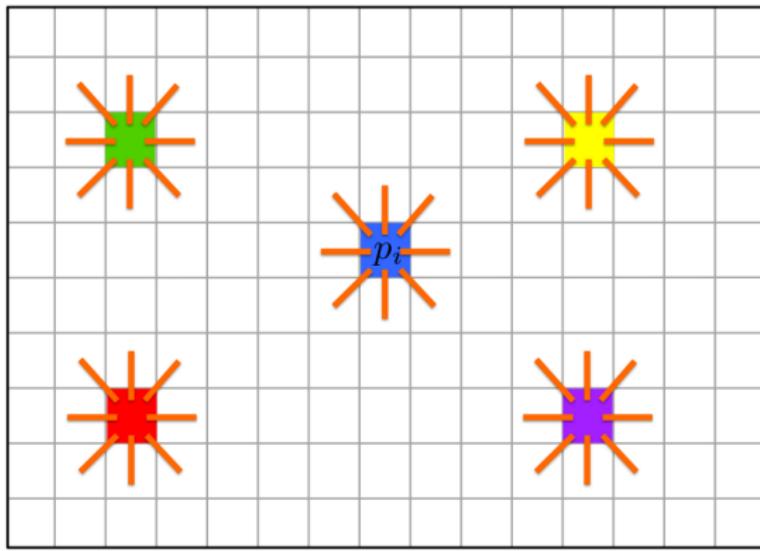
$$I \quad w_{ij} = |I(p_i) - I(p_j)|$$



Felzenswalb & Huttenlocher

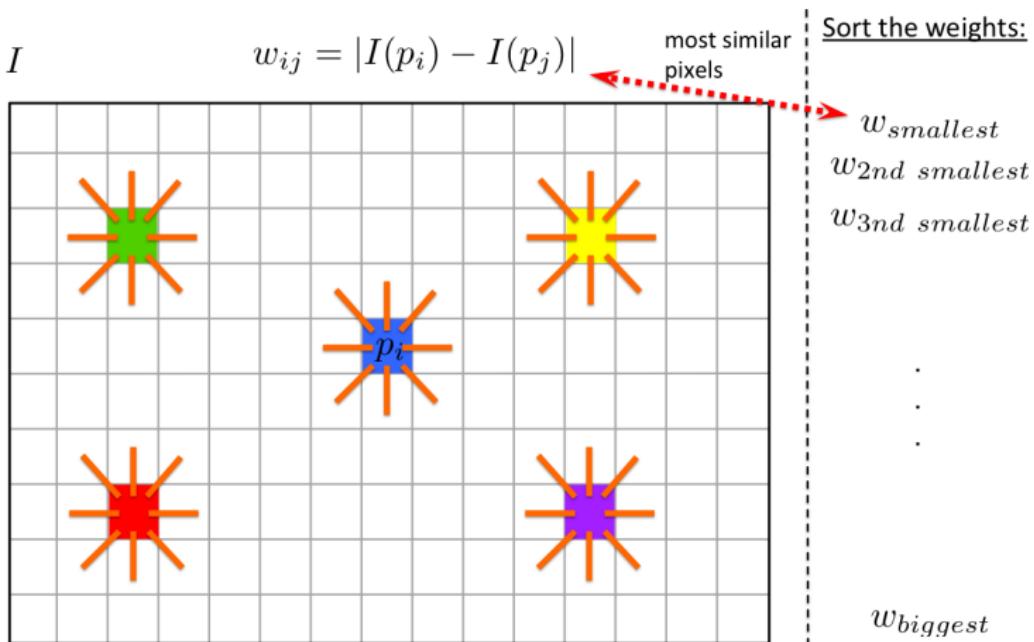
- And do that for **all pixels** in the image.

$$I \quad w_{ij} = |I(p_i) - I(p_j)|$$



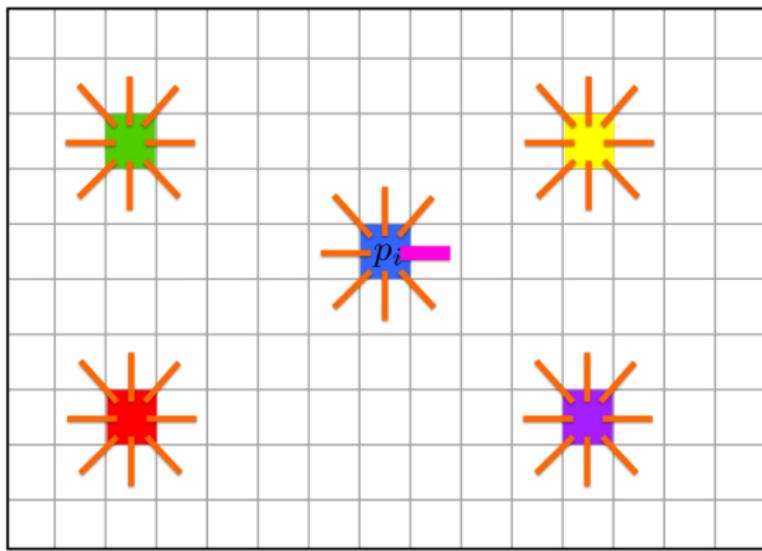
Felzenswalb & Huttenlocher

- Now sort all the weights in the image in ascending order.



Felzenswalb & Huttenlocher

- Pick the weight from the top of the list (most similar two neighboring pixels)

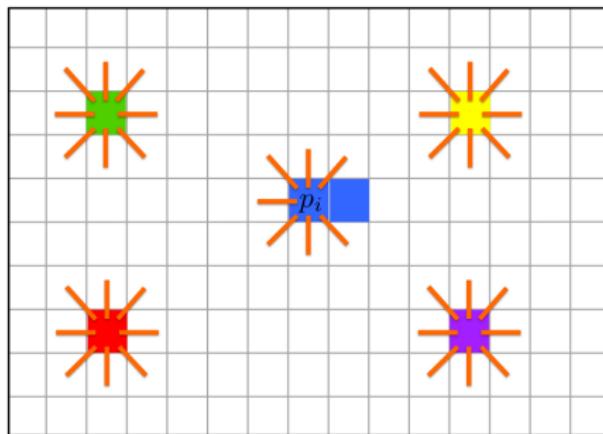


Sort the weights:

→ $w_{smallest}$
 $w_{2nd\ smallest}$
 $w_{3nd\ smallest}$
.
. .
 $w_{biggest}$

Felzenswalb & Huttenlocher

- If the weight is smaller than the internal dissimilarities (plus a threshold), **merge** the pixels
- We define the internal dissimilarities as the largest difference in the minimal spanning tree of a component. i.e. we check if a pixel (or component) is at least as similar to the group it is being added to, as the group is with itself.



Sort the weights:

→ $w_{smallest}$

$w_{2nd\ smallest}$

$w_{3nd\ smallest}$

.

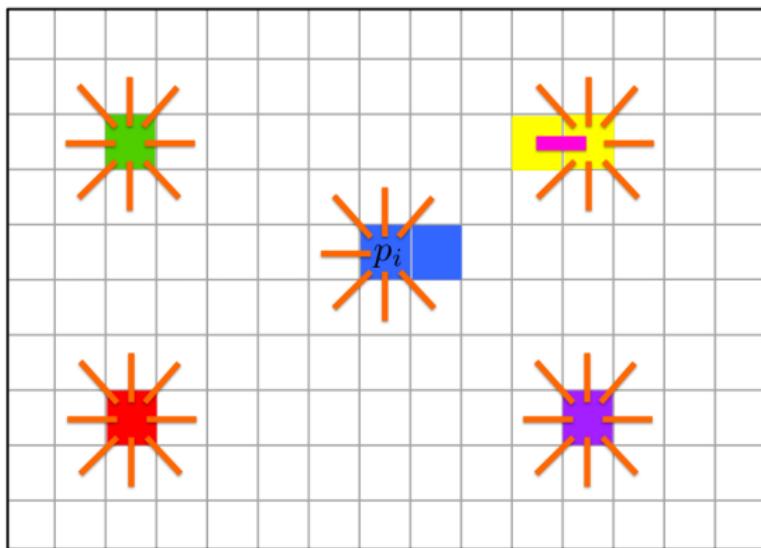
.

.

$w_{biggest}$

Felzenswalb & Huttenlocher

- Take the second weight in the list and do the same.



Sort the weights:

$w_{smallest}$

$w_{2nd\ smallest}$

$w_{3nd\ smallest}$

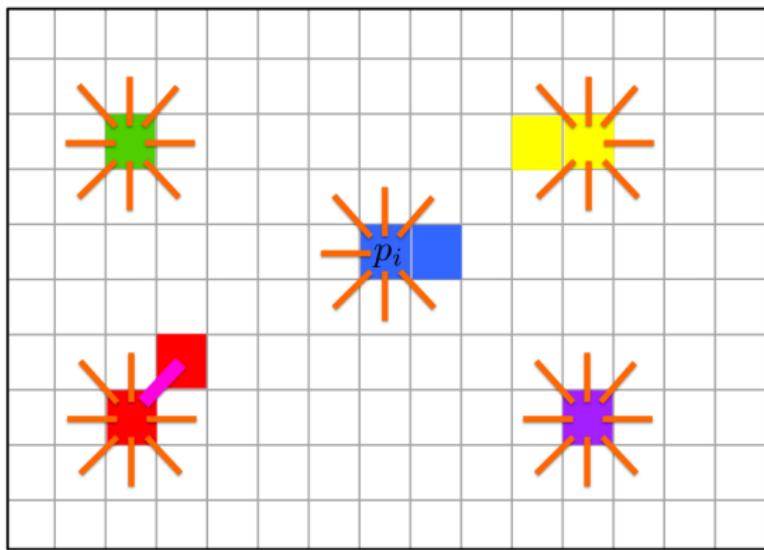
.

.

$w_{biggest}$

Felzenswalb & Huttenlocher

- And the third one. Repeat until the end of list, or until the weight is higher than the dissimilarity between the pair of components.



Sort the weights:

$w_{smallest}$

$w_{2nd\ smallest}$

→ $w_{3nd\ smallest}$

.

.

.

$w_{biggest}$

Felzenswalb & Huttenlocher

- Result. The algorithm runs real-time.



Segmentation Algorithms

There are a lot of them out there

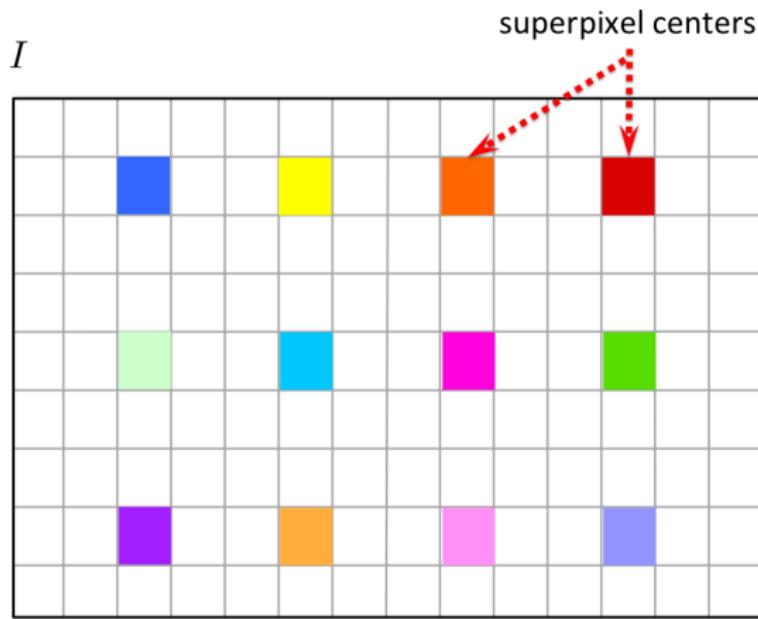
- Felzenswalb and Huttenlocher's Graph-based Segmentation
- SLIC → Let's quickly look at this one too
- SEEDS
- Shi and Malik's Normalized Cuts
- Malik's group: Probability of boundary (gPb)
- Grundmann et al: Hierarchical Graph-based Video Segmentation
- Chang et al., Temporal Superpixels

For most the code is available!

SLIC: <http://ivrg.epfl.ch/research/superpixels>

SLIC Superpixels

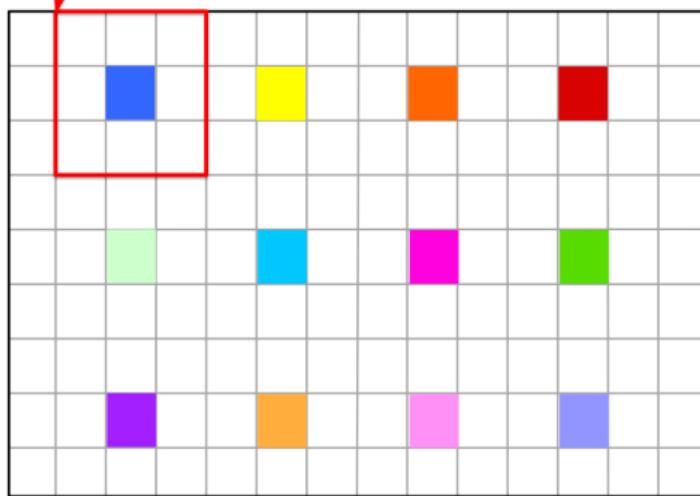
- Ask the user for the number of superpixels.
- Initialize the superpixel **centers** in a regular grid.



SLIC Superpixels

- Allow each center to shift in a 3×3 window to the location with the smallest gradient.

Shift each center to position with smallest gradient in
3x3 patch around center. Any idea why this is a good
strategy (hack)?
 I

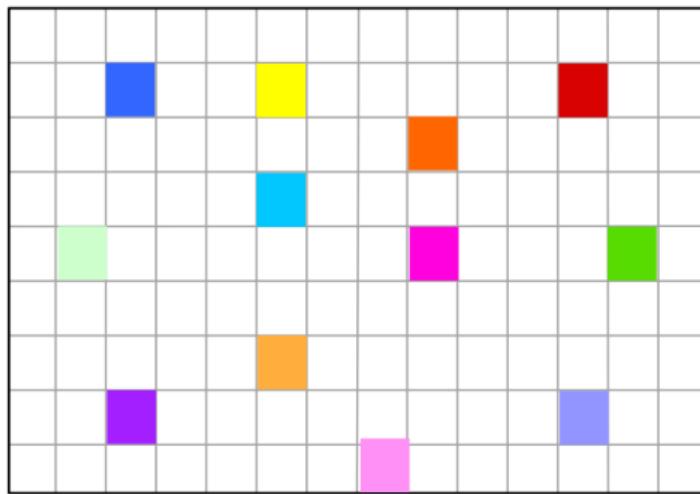


SLIC Superpixels

- Compute feature vectors for each pixel.

Compute a feature vector in each pixel (and superpixel center). Feature vector: [color, x, y]

I

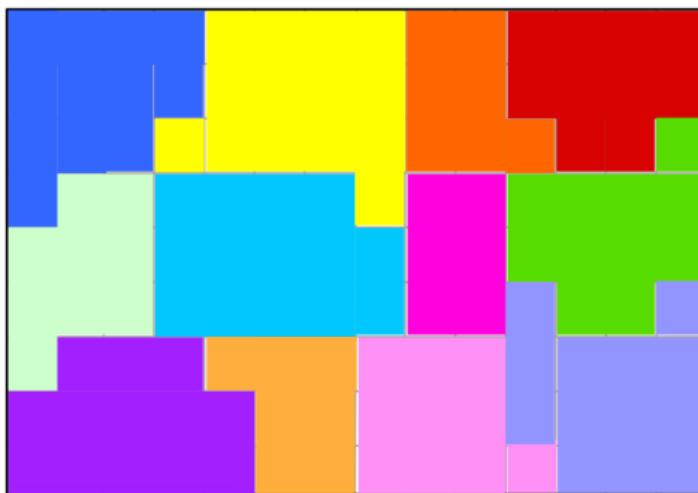


SLIC Superpixels

- Assign each pixel to the closest superpixel center.

Assign all other pixels to the closest center, where ``closest'' is looking at a distance between the feature vectors

I

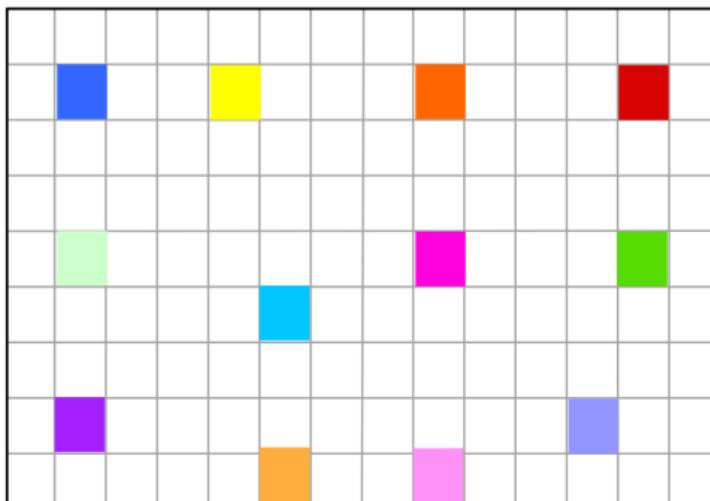


SLIC Superpixels

- Re-compute the centers (just average the features across pixels in each superpixel).

Re-compute the superpixel centers (just compute the average feature vector).

I Also compute how much each center moved from previous iteration.

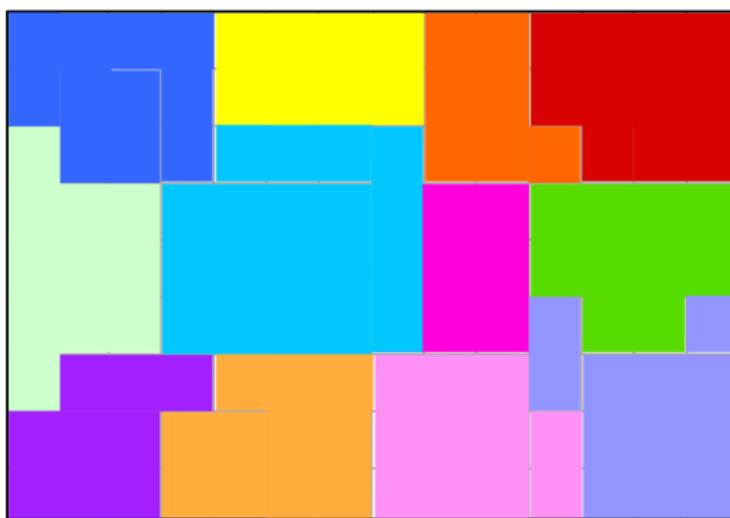


SLIC Superpixels

- Assign again. Iterate until the centers become stable (don't move from the previous iteration).

Assign again, and iterate. Stop when none of the centers moves more than a threshold.

I



Results

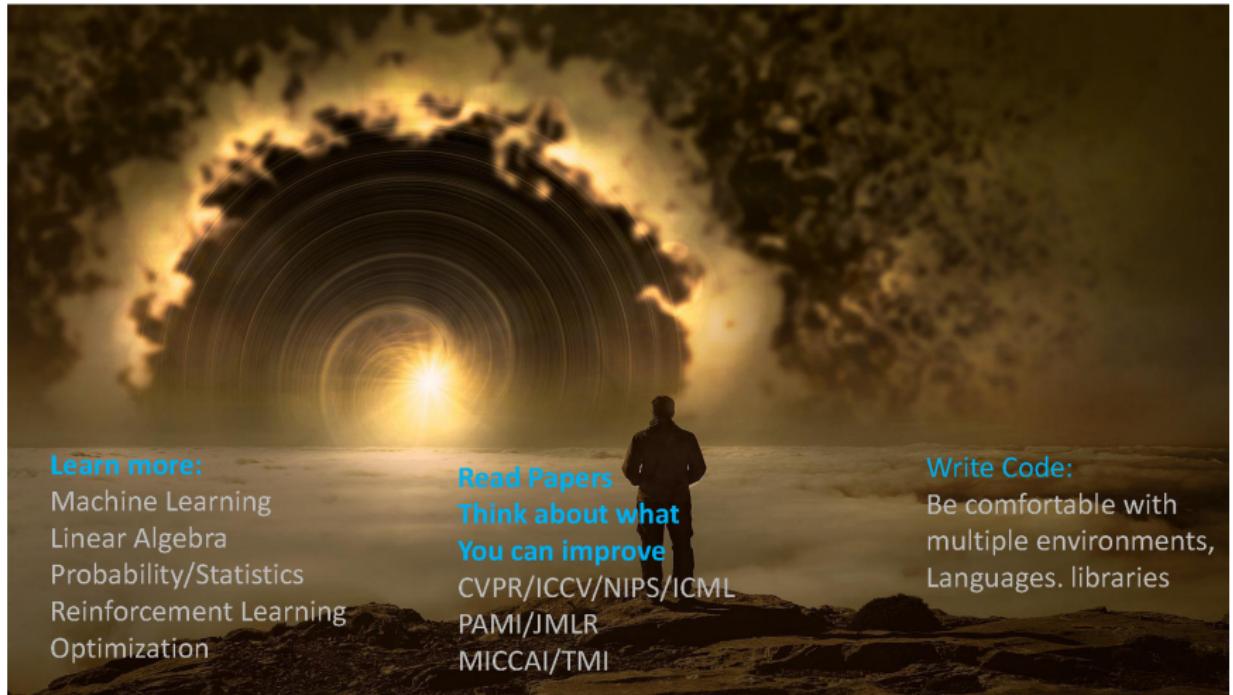
- Results



Is there Life After CSC420?



Is there Life After CSC420?



Learn more:

Machine Learning
Linear Algebra
Probability/Statistics
Reinforcement Learning
Optimization

Read Papers

Think about what
You can improve
CVPR/ICCV/NIPS/ICML
PAMI/JMLR
MICCAI/TMI

Write Code:

Be comfortable with
multiple environments,
Languages. libraries

Is there Life After CSC420?



CSC2548 Machine Learning in Computer Vision

CSC2305 Numerical Methods for Optimization

CSC2506 Probabilistic Learning and Reasoning

CSC2511 Natural Language Computing

CSC2503 Foundations of Computer Vision

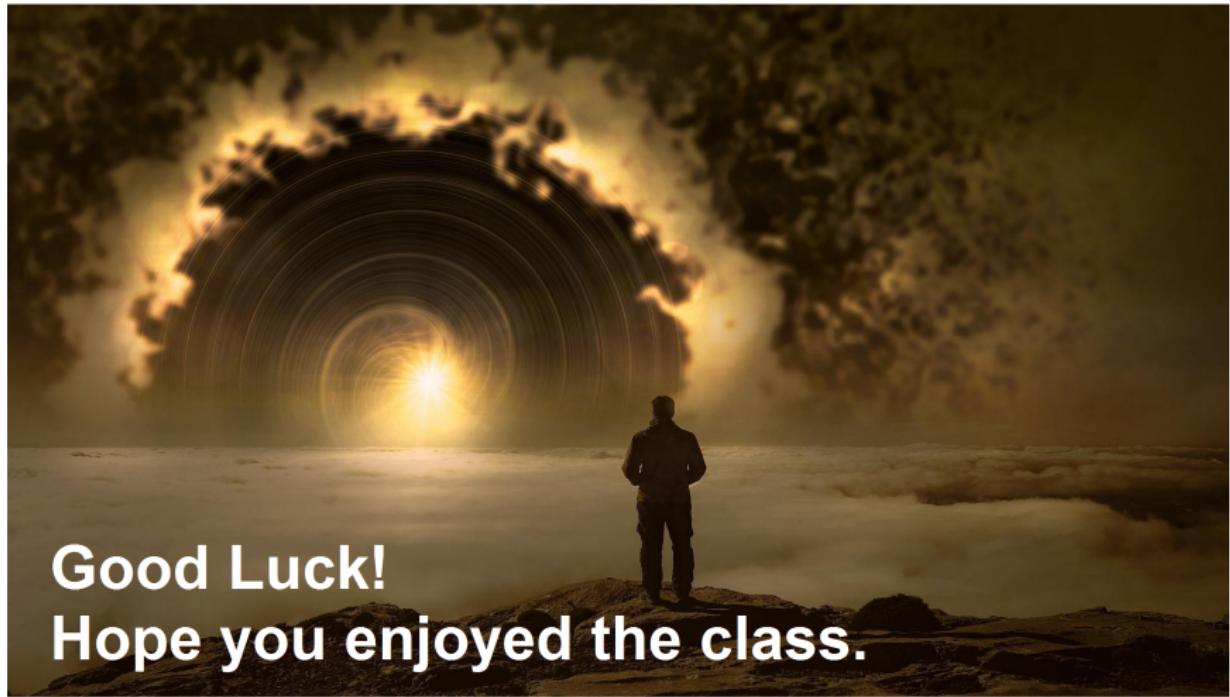
CSC2541 Topics in Machine Learning: Uncertainty

CSC2523 Object Modeling: Shape Perception

CSC3547 Current Algorithms in Machine Learning

CSC2548 Machine Learning in Computer Vision

Is there Life After CSC420?



**Good Luck!
Hope you enjoyed the class.**